


Exp 3 — Ablação de quantização: bf16 vs 8-bit

Qwen2.5-7B



Llama-3.1-8B

